

DRISHTI CHADHA

Product Manager — Enterprise Data & AI Platforms

Bengaluru, India • +91 7579681146 • drishtichadhaa@gmail.com • linkedin.com/in/drishtichadha

PROFESSIONAL SUMMARY

Technical Product Manager with 5+ years driving **enterprise data platforms** from ambiguous business need through PRDs, architecture, launch, and adoption — personally owning product strategy and delivery for **35+ enterprise Databricks/AI accounts** at Aays / Aidetic. Direct, hands-on experience **unifying fragmented enterprise data**: personally migrated 7+ organizations off siloed Redshift, Glue, and BI stacks onto governed Lakehouse architectures with **Unity Catalog** as the security and governance layer, and single-handedly built and shipped **AutoDBX**, an automation framework that cut migration time **60% across 7+ engagements**. Shipped **four production AI experiences 0-to-1** (car-damage detection, NLP resume-matching, distributed training/vector search, MLOps monitoring). Writes detailed product specs, drives Go/No-Go decisions and executive sign-off across matrixed, geographically distributed stakeholder groups — including traveling from India to the UK to lead UAT and secure executive sign-off on a global supply-chain platform — and uses adoption and quality telemetry to define success metrics and prioritize the roadmap. **Databricks Certified Data Engineer (Associate & Professional)**. B.Sc. Physics, University of Delhi.

RELEVANT STRENGTHS FOR THIS ROLE

- **Unifying fragmented enterprise data**: Personally led migrations that replaced Redshift, Glue ETL, Superset, and QlikView across 7+ enterprises with governed Delta Live Tables / Databricks SQL pipelines on Unity Catalog — the same fragmentation-to-unified-platform problem this role owns.
- **Built automation that scales a platform**: Single-handedly designed and shipped AutoDBX, a Databricks Asset Bundle automation framework, cutting migration time 60% across 7+ engagements — now a reusable accelerator adopted firm-wide, not just a one-off project.
- **Project Thema (Global CPG, London)**: Architected a greenfield Medallion Lakehouse spanning Customer Service Level, Demand Forecasting, Warehouse, and Inventory Management, integrating SAP, BlueYonder, and manual sources on Databricks and Azure. Shipped Phase 1 in 8 weeks and a global Phase 2 rollout in 8 weeks to ~500 stakeholders; personally traveled from India to the UK to lead UAT and secure executive sign-off, eliminating ~25,000 hours/year of manual reporting, driving 1,500+ hours of additional time savings within the year, and unlocking ~\$5M in avoided inventory write-offs.
- **PRDs, specs, and governed delivery**: Authored product specs and system designs end-to-end (discovery through go-live) across 35+ accounts, personally securing executive sign-off at each — with Unity Catalog governance built in for auditable, repeatable deployment.
- **Cross-functional, geographically distributed delivery**: Sole point of contact across Databricks field engineering, customer technical teams, and account stakeholders spanning global clients in the US, UK, and Australia; personally ran an 8-week global rollout to ~500 stakeholders across a distributed org.
- **Metrics-driven iteration**: Used adoption telemetry and model-monitoring data to define platform success metrics and clear post-launch blockers — not just technical go-live.

EXPERIENCE

Aays / Aidetic | Bengaluru, India

Jun 2025 – Aug 2026

Assistant Manager – AI & Data Platform Architecture (Databricks Practice)

- **Owned product strategy and technical delivery for 35+ enterprise accounts**: designed Lakehouse architectures on Unity Catalog, Delta Live Tables, and Databricks Asset Bundles, personally translating business requirements into product specs and system design end-to-end — discovery through go-live — securing executive sign-off at each account.
- **Unifying fragmented enterprise data**: personally architected and led migrations for Scripbox, Urbint, and IBS Software (AWS Redshift to Databricks — schema redesign, pipeline re-engineering, Genie integration); Lenden and Phamarack (AWS Glue ETL to Delta Live Tables); Clickpost (legacy ETL to Delta Live Tables); and Myntra (Superset to Databricks SQL) — each landed on Unity Catalog governance for auditable, repeatable deployment.
- **Built AutoDBX**, a Databricks Asset Bundle automation framework, single-handedly, cutting migration time 60% across 7+ engagements — plus Business-Driven MLOps and GPU-Optimisation accelerators, both adopted firm-wide.
- **AI products shipped 0-to-1**: a car-damage detection model for Acko Insurance; an NLP JD-and-resume matching model for TurboHire; a distributed-training and vector-search system for Mihup AI; and an end-to-end MLOps monitoring/retraining pipeline for Observio.

- Served as sole point of contact across a matrixed, geographically distributed org — Databricks field engineering, customer technical teams, and account stakeholders — personally resolving integration and architecture escalations across joint reviews spanning all 35+ accounts.
- Ran production launches and post-go-live monitoring, hardening pipelines and models for reliability at scale, and tracking adoption metrics to clear post-launch blockers.
- Helped drive Aays / Aidetic to Databricks Silver Partner & DPP status in under 10 months through certified-practitioner enablement and production delivery.

Aays / Aidetic | Bengaluru, India

Mar 2024 – Jul 2025

Functional Consultant & Project Manager (Technical Architecture & Delivery)

- Led technical implementation for 15+ concurrent Fortune 500 engagements across Supply Chain, Finance, ESG, Sales, and Health — designing solution architecture from requirements through executive sign-off.
- **Project Thema (Global CPG, London):** architected a greenfield Medallion Lakehouse spanning Customer Service Level, Demand Forecasting, Warehouse, and Inventory Management, integrating SAP, BlueYonder, and manual sources on Databricks and Azure. Shipped Phase 1 in 8 weeks and a global Phase 2 rollout in 8 weeks to ~500 stakeholders; personally traveled from India to the UK to lead UAT and secure executive sign-off, eliminating ~25,000 hours/year of manual reporting, driving 1,500+ hours of additional time savings within the year, and unlocking ~\$5M in avoided inventory write-offs.
- Architected and delivered the first wave of Lakehouse migrations across 7+ enterprises — replacing Redshift, Glue ETL, Superset, and QlikView with Delta Live Tables and Databricks SQL — achieving up to 50% cost reduction.

Aays / Aidetic | Gurugram, India

Mar 2022 – Mar 2024

Senior Data Analyst / Data Analyst

- Built ML/AI forecasting systems on Azure for a USD 40Bn CPG conglomerate (Mars Food): Daily Business Position forecasting (FB Prophet + ARIMA) and anomaly detection across 20+ external and 3 internal factors at <5% MAPE; delivered 60+ user-centric Power BI dashboards.
- Designed Azure Data Factory pipelines and built a scalable WHO Nutrition & Health reporting data model integrating 5 source systems, automating regulatory reporting and eliminating 90% of manual collation effort.
- Built a unified SQL data model and FbProphet forecasting system for an Australian fashion brand covering 30+ sales and inventory reports, including SKU-return prediction.

Namekart | India

May 2020 – Oct 2021

Business Analyst

- Grew clientele 20%, led a team of 5, and supported \$1M+ in customer acquisitions through analytics-led engagement.

SKILLS, EDUCATION & CERTIFICATIONS

Product & Program Management: PRD & product-spec authoring (inputs, outputs/UI, processing logic), Go/No-Go decision framing & executive sign-off, cross-functional/matrixed program delivery, roadmap prioritization, adoption & telemetry-driven iteration, stakeholder management

Data Platform & Governance: Lakehouse/Medallion architecture design, Unity Catalog governance design, legacy-system unification (SAP, BlueYonder, Redshift, Glue), technical scoping & solution design, production hardening & reliability engineering

AI Product Delivery: MLflow model monitoring, LLM & agent evaluation, NLP, distributed training & vector search, XGBoost, LightGBM, PyTorch, TensorFlow, FB Prophet, ARIMA/SARIMA

Data & Cloud Platforms: Databricks (Unity Catalog, Delta Live Tables, Delta Lake, Lakeflow, Asset Bundles, Genie, Databricks SQL, PySpark, Spark SQL), Azure (ADF, ADLS, Purview), AWS (Redshift, Glue)

Languages & Tools: Python, SQL, PySpark | Power BI, QlikView, Superset, Streamlit

Education: B.Sc. Physics, Hansraj College – University of Delhi (2017–2020), CGPA 8.09, Top 10% of Class

Certifications: Databricks Certified Data Engineer Associate • Databricks Certified Data Engineer Professional • Databricks Partner Sales Ready (Mar 2026 – Mar 2028) • PL-300 Microsoft Certified Data Analyst Associate • SQL for Data Science (UC Davis) • IBM Data Analysis with Python